Hemanth Vallamkonda

Five years in finance and operations; builds governed AI workflows • Right to work in the UK, no sponsorship • vallamkondahemanth@gmail.com • +44 7553 909828

FINANCE OPERATIONS LEAD (AI) · FOR A GLOBAL FINTECH, VIA CLIENT SERVER

AI spend, made accountable

The role's first month, on a made-up FinTech on Azure: AI spend metered, reconciled, owned and decided.

Illustrative usage; published list prices read 3 and 4 November 2026 (sources (https://growthforge.github.io/applications/client-server-ai-spend-ledger/build.html#sources)).

October decision memo

ILLUSTRATIVE: A MADE-UP ORGANISATION

TO Finance director, use-case owners FROM Hemanth Vallamkonda (illustrative memo)

AI spend landed at $21,215, 107% of its $19,770 budget. The budget is 115% of the August and September average **S3** (https://growthforge.github.io/applications/client-server-ai-spend-ledger/build.html#src-S3); spend was 123% of that average. The ledger's forecast alert fired 13 October; actual spend crossed the budget 30 October.

WHY IT MOVED

September to October, AI spend (+$4,032)

September	**$17,183**	
Volume	**+$1,390**	per day; mostly KYC document review, runs a day up 34%
Calendar	**+$562**	31 days against 30, at September's daily level
Unit price	**+$286**	`gpt-support-hard` moved from GPT-5.2 to GPT-5.4 on 8 October, at list price
Mix and consumption	**+$109**	largest in Support assistant: tokens and billed cost per run
Fixed shared platform	**+$8**	AI Search and the AI gateway, fixed daily charges; beyond the extra day
Unattributed AI	**+$1,678**	billed AI with no use case: `gpt54-sandbox` +$1,356, `foundry-prod-uks` +$321
October	**$21,215**	

CONTROL CHECKS

91% of spend has a named owner (Run; Walk 85% **S5** (https://growthforge.github.io/applications/client-server-ai-spend-ledger/build.html#src-S5)); `log-shared-uks` is untagged and $558 of unmetered AI use is held.

Gateway metering explains 85% of token-billed AI cost; `gpt54-sandbox` bypasses the gateway.

Anomaly, 14 to 16 Oct: KYC document review failures rose to 38% of runs from 11%, retries 2.4× usual (illustrative cause: a retry storm on long documents). The retry cap is a guardrail, not a saving.

DECISIONS

▲ SCALE

Support assistant

$0.025 avoidable per accepted draft
value $2.40
84% succeed (bar 80%)

Value 97× avoidable cost, bar met.

▲ SCALE

Fraud alert explanations

$0.0048 avoidable per explanation used
value $0.60
46% succeed (bar 40%)

Value 126× avoidable cost, bar met.

○ FIX

KYC document review

$0.19 avoidable per approved summary
value $12.00
66% succeed (bar 85%)

34% of cost went on failed or rejected runs. Next: pre-check documents and cap retries on long PDFs (guardrail).

○ FIX

Internal knowledge assistant

$0.075 avoidable per useful answer
value $0.55
$0.60 with fixed shared cost
39% succeed (bar 60%)

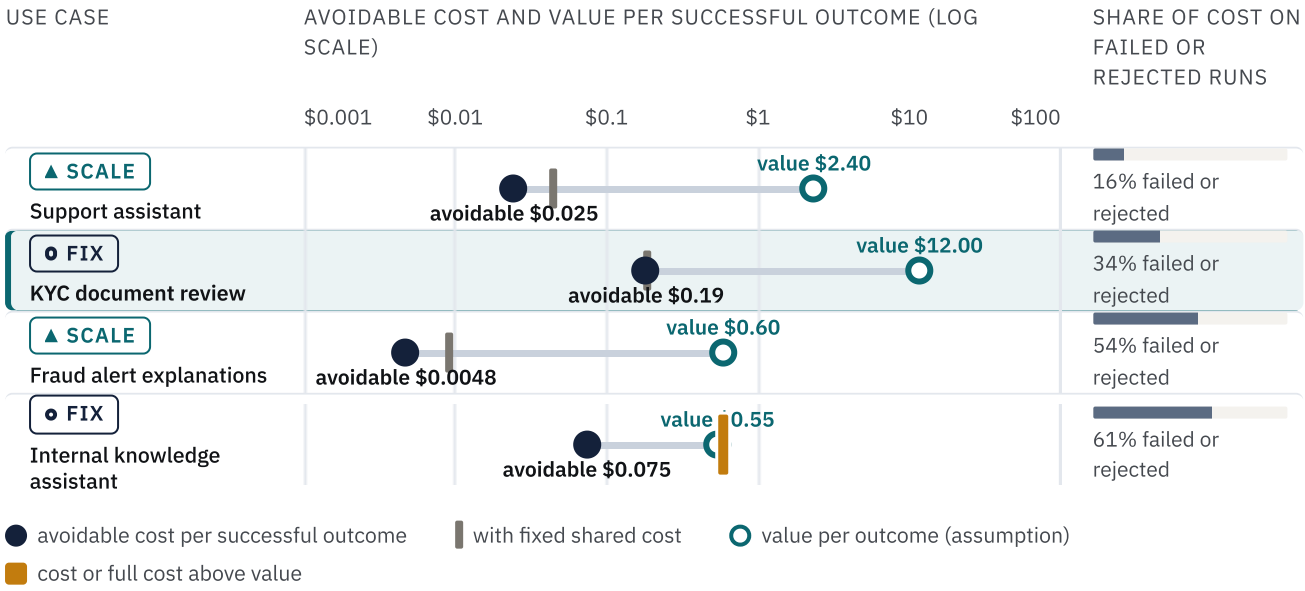
Full cost is above value. Fix retrieval and downsize the tier: $0.10 fully loaded.

Unit economics

Avoidable cost against value per successful outcome (assumptions in value.json (https://growthforge.github.io/applications/client-server-ai-spend-ledger/pack/data/value.json)), and the share lost to failures.

ILLUSTRATIVE

ILLUSTRATIVE Avoidable cost per successful outcome against its value, October. Log scale: each gridline is ten times the last. Decisions are for October.



KYC document review, October: P50 $0.17 and P90 $0.18 of daily avoidable cost per approved summary; 34% of cost on failed or rejected runs; its share of fixed shared cost (AI Search and the AI gateway) is $86, 2% of allocated cost, which stays if it stops; cache hit rate 10%; 84k tokens per approved summary. The 14 to 16 Oct incident days ran at $0.26 to $0.29 per approved summary against $0.17 normally, and explain 93% of the month's excess failures; they lift the monthly figure above the daily P90.

Unit economics, October. Illustrative: a made-up organisation. Avoidable cost is what stops if the use case stops (model tokens and embeddings); every run counts, including failed, retried and discarded ones. The daily range is of avoidable cost.

USE CASE	ALLOCATED AI SPEND	SUCCESSFUL OUTCOMES	AVOIDABLE COST PER OUTCOME	WITH FIXED SHARED COST	VALUE PER OUTCOME	FAILED OR REJECTED
Support assistant Customer Operations	$6,131	137,974 accepted drafts	**$0.025**	$0.044	$2.40	16%
KYC document review Financial Crime	$5,546	29,501 approved summaries	**$0.19**	$0.19	$12.00	34%
Fraud alert explanations Fraud Analytics	$3,180	341,395 explanations used	**$0.0048**	$0.0093	$0.60	54%
Internal knowledge assistant Engineering	$4,353	7,305 useful answers	**$0.075**	$0.60	$0.55	61%
Use-case AI spend	**$19,210**	includes $8,159 of fixed shared platform (AI Search by retrieval calls, the AI gateway by requests); unattributed AI (no use case: the sandbox and unmetered use) $2,006; all AI spend $21,215				

Savings pipeline

Every saving shows its method and evidence grade; levers apply in sequence; the next bill checks unit cost. ILLUSTRATIVE

MEASURED
$9,545 a month

WITH ASSUMPTION-BASED
$17,043 a month

WITH FOUNDATION PRIORS
$17,232 a month

REALISED IN OCTOBER
$2,302 in the month

Monthly savings by lever, priced from the October ledger. Illustrative.

LEVER	SAVING A MONTH	EVIDENCE	METHOD AND QUALITY GATE	NEXT-BILL CHECK	STATUS
▲ SCALE **Support assistant**					
Route simple hard-case tickets from GPT-5.4 to GPT-5.4 mini Owner: Customer Operations	$444	assumption-based	Assumption-based: 40% of hard-case tickets are simple enough (assumption); the price sheet puts the same tokens on GPT-5.4 mini at 30% of GPT-5.4, so routing saves 70% of each routed ticket's cost. Gate: Acceptance on routed tickets within 2 points of GPT-5.4 in a 200-ticket shadow test Routing starts a new prompt cache on the mini deployment, so the caching lever is priced after it.	avoidable cost per accepted draft: $0.025 now, $0.022 or less next bill	identified
Put the stable 4,000-token prefix first so automatic caching hits Owner: Customer Operations	$148	assumption-based	Assumption-based: hit rate 17.4% measured; the stable 4,000-token prefix is 24.2% of the average 16,508-token prompt, so at most 24.2% of input can hit. Cached input at 10% of the input price then saves 7.2% of input cost. Gate: No change to draft content: the prompt is reordered, not rewritten	avoidable cost per accepted draft: $0.025 now, $0.020 or less next bill	identified

LEVER	SAVING A MONTH	EVIDENCE	METHOD AND QUALITY GATE	NEXT-BILL CHECK	STATUS
Cap reply length and ask for structured drafts Owner: Customer Operations	$93	☐ Foundation prior	Foundation prior: 10% to 40% range, counted at 10%. Gate: Agent acceptance rate does not fall	avoidable cost per accepted draft: $0.025 now, $0.020 or less next bill	identified
FIX KYC document review					
Batch the overnight KYC document runs Owner: Financial Crime	not counted	not graded	Blocked: Anthropic's batch platforms list does not include Microsoft Foundry. • Run the overnight backlog through a platform that lists batch for Claude (the Claude API), after a procurement and data-residency review • Stay on Foundry and cache the shared review instructions instead Gate: Not applicable until a platform route is agreed	none until priced	blocked
Pre-check documents and cap retries on long PDFs (guardrail) Owner: Financial Crime	$265 once, not counted	■ measured	Guardrail, not a recurring saving: the 14 to 16 October incident cost $265 in failed runs above the normal rate (measured). A page-count pre-check and a two-retry cap stop a repeat. Gate: Approval rate on documents that pass the pre-check does not fall	none until priced	guardrail
▲ SCALE Fraud alert explanations					
Generate explanations only when an analyst opens the alert Owner: Fraud Analytics	$675	◧ assumption-based	Assumption-based: 51.8% of explanations were never used (measured); generating on open avoids them, assuming 80% of unused ones are on alerts nobody opens. Gate: Explanation shown within a few seconds of the analyst opening the alert	avoidable cost per explanation used: $0.0048 now, $0.0028 or less next bill	identified

LEVER	SAVING A MONTH	EVIDENCE	METHOD AND QUALITY GATE	NEXT-BILL CHECK	STATUS
Send the overnight explanation backlog through the Batch API Owner: Fraud Analytics	$191	assumption-based	Assumption-based: 40% of explanations are for the overnight re-scoring backlog and can wait up to 24 hours (assumption from the team's process; confirm from gateway timestamps); Batch is 50% off input and output on Azure Global Standard. Gate: Backlog explanations ready before the morning review	avoidable cost per explanation used: $0.0048 now, $0.0022 or less next bill	identified
FIX **Internal knowledge assistant**					
Trim retrieved context to the top passages Owner: Engineering	$96	Foundation prior	Foundation prior: 20% to 60% range, counted at 20%. Gate: Share of answers rated useful does not fall	avoidable cost per useful answer: $0.075 now, $0.062 or less next bill	identified
Shared AI platform					
Right-size the shared AI Search tier Owner: Engineering	$5,250	assumption-based	Assumption-based: 6 Standard S2 units today; 3 Standard S1 units ($0.336 an hour each against $1.34, Azure Retail Prices API) if the index and query load fit, which is the assumption to test. Gate: Index size and query latency within S1 limits for both users before the move	search tier cost per day: $194 now, $24.19 or less next bill	identified
Cloud estate					
Rightsize VMs that stayed under the CPU threshold all quarter Owner: Fraud Analytics, Financial Crime	$496	assumption-based	Assumption-based: 2 VMs (vm-fraud-scoring-01, vm-kyc-ocr-01) stayed under 20% P95 CPU all quarter (illustrative telemetry); one size down is priced from the Azure Retail Prices API, which halves the Dsv5 rate. Memory must fit too. Gate: P95 CPU on the new size stays under the owner's limit for two weeks	cost per VM-day: $15.98 now, $7.99 or less next bill	identified

LEVER	SAVING A MONTH	EVIDENCE	METHOD AND QUALITY GATE	NEXT-BILL CHECK	STATUS
Scale non-production AKS to zero at night and at weekends Owner: Engineering	$3,755	measured	Measured from the bill: mean daily cost $213 in September, $92 from the change; 64% of hours are scheduled off. Gate: Test environments up by 07:00 on working days	non-production cluster cost per day: $213 now, $91.98 or less next bill	realised
Savings plan for the steady compute floor Owner: Engineering	$5,790	measured	Measured floor, price-sheet discount: production compute never fell below $751 a day at pay-as-you-go rates ($31.30 an hour) across 92 days, after the rightsizing above. A 1-year savings plan bills Dsv5 at 75.1% of pay-as-you-go (Azure Retail Prices API), so commit $23.52 an hour: the floor times (1 – 24.9%). Gate: Rightsizing done first; the commitment is reviewed when the floor moves	effective cost per production compute hour at the floor: $31.30 now, $23.52 or less next bill	identified
Lifecycle tiering for logs and cost exports Owner: Engineering	$296	assumption-based	Assumption-based: 70% of the data is older than 30 days (assumption); Cool LRS at $0.0105 against Hot at $0.0192 per GB-month (Azure Retail Prices API). Cool's higher operation charges and early-deletion charge are not priced. Gate: Audit retention periods still met	storage cost per GB-month: $0.019 now, $0.013 or less next bill	identified
Provisioned throughput (PTU) for AI Owner: Engineering	not priced	not graded	Not yet. Daily production tokens swing 2.3× between the quietest and busiest day and AI volume is still moving; PTU sizing needs the Foundry capacity calculator, so it is not priced. Gate: a stable 60 to 90 day token baseline	none	not yet

Background that fits

Reporting packs.
At Kuruvella & Co., ~35% less report preparation time across 30+ SME client entities.

Measuring a unit of work.
At Bnk To The Future, ~25% faster turnaround across 100+ cases, with risk-committee reporting.

Governed review steps.
At MTest Labs, test design went from an 8-hour day, or 2 to 3 days, to 1 to 3 supervised hours; a tester approves every case.

Contact

How it works (https://growthforge.github.io/applications/client-server-ai-spend-ledger/build.html) has the code; the pack (https://growthforge.github.io/applications/client-server-ai-spend-ledger/pack.zip) runs with `node --test`.

vallamkondahemanth@gmail.com **+44 7553 909828**

Prepared by Hemanth Vallamkonda as part of an application to a global FinTech via Client Server. Independent and illustrative; uses public information only.

The organisation, its usage and its bill are made up. Not affiliated with or endorsed by Client Server or its client. Prices: Anthropic and Azure Retail Prices list prices, read 3 and 4 November 2026.

vallamkondahemanth@gmail.com · +44 7553 909828 · Right to work in the UK; no sponsorship required.